

imtiaaz-mall-eda

March 12, 2024

```
[33]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd
from sklearn.cluster import KMeans, DBSCAN, kmeans_plusplus
```

```
[34]: #open the data file
df = pd.read_json('electronics.json')
df.head()
```

```
[34]:
```

	Customer_ID	Age	Gender	Income_Level	\
0	b81ee6c9-2ae4-48a7-b283-220eaa244f43	40	Female	Medium	
1		25	Male	High	
2	fdf79bcd-5908-4c90-8501-570ffb5b7648	57	Other	Low	
3	878dccba-893a-48f9-8d34-6ed394fa3c9c	38	Female	Medium	
4	0af0bd81-73cc-494e-aa5e-75c6d0b6d743	68	Other	Medium	

	Address	\
0	43548 Murray Islands Suite 974\nAmyberg, CT 13457	
1		
2	79683 Kevin Hill Apt. 555\nJohnshire, AR 39961	
3	02998 Hall Meadows Suite 809\nNorth Robertvill...	
4	21411 Timothy Ford Apt. 320\nDavisborough, AR ...	

	Transaction_ID	Purchase_Date	\
0	c6a6c712-e36b-406a-bfde-f53bdcf4744f	2022-04-26	
1	0b587838-1e4f-4231-b488-42bcd47c052a	2021-08-10	
2	462925b1-a5bf-4996-bda2-59749de64eea	2021-12-09	
3	3cfafa02-6b34-4d77-9e05-d223dfab64e8	2022-12-03	
4	0d8dc27a-0c8f-4a82-b57e-8bf54cee9759	2020-06-08	

	Product_ID	Product_Category	Brand	\
0	d2f767d6-b01a-41a2-87f7-ec1d1186f50e	Clothing	Brand_C	
1	79eadc55-2de1-41cf-b1b6-40118c0bf8ec	Books	Brand_A	
2	9ab75a68-4329-4bd9-a259-2233c0f34c93	Electronics	Brand_A	
3	d518569b-ff79-494b-b2b6-7e2af39db86a	Clothing	Brand_C	
4	b6deac9d-2b7e-4a51-8273-a6534910b3bc	Books	Brand_B	

	Purchase_Amount	Average_Spending_Per_Purchase	Purchase_Frequency_Per_Month	\
0	193	59		2
1	318	77		2
2	197	100		9
3	262	97		3
4	429	85		7

	Brand_Affinity_Score	Product_Category_Preferences	Month	Year	Season
0	2	Low	01	2010	Winter
1	1	Low	08	1989	Fall
2	1	Low		1995	Winter
3	4	Low	09	2012	Fall
4	2	High	01	2010	Summer

```
[35]: df.replace('', np.nan).isna().sum()
```

```
[35]: Customer_ID      32
      Age              33
      Gender           33
      Income_Level     41
      Address          32
      Transaction_ID   39
      Purchase_Date    35
      Product_ID       40
      Product_Category 44
      Brand            46
      Purchase_Amount  33
      Average_Spending_Per_Purchase 26
      Purchase_Frequency_Per_Month 37
      Brand_Affinity_Score 47
      Product_Category_Preferences 31
      Month            40
      Year             39
      Season           36
      dtype: int64
```

```
[36]: df.replace('Hidden', np.nan).isna().sum()
```

```
[36]: Customer_ID      12
      Age              7
      Gender           15
      Income_Level     9
      Address          15
      Transaction_ID   11
      Purchase_Date    13
      Product_ID       9
```

Product_Category	16
Brand	12
Purchase_Amount	16
Average_Spending_Per_Purchase	14
Purchase_Frequency_Per_Month	18
Brand_Affinity_Score	14
Product_Category_Preferences	12
Month	13
Year	13
Season	12
dtype: int64	

```
[37]: df.replace('', np.nan, inplace=True)
df.replace('Hidden', np.nan, inplace=True)

df.isna().sum()
```

[37]: Customer_ID	44
Age	40
Gender	48
Income_Level	50
Address	47
Transaction_ID	50
Purchase_Date	48
Product_ID	49
Product_Category	60
Brand	58
Purchase_Amount	49
Average_Spending_Per_Purchase	40
Purchase_Frequency_Per_Month	55
Brand_Affinity_Score	61
Product_Category_Preferences	43
Month	53
Year	52
Season	48
dtype: int64	

```
[38]: #Convert column to numeric
df['Age'] = pd.to_numeric(df['Age'], errors='coerce')
df['Purchase_Amount'] = pd.to_numeric(df['Purchase_Amount'], errors='coerce')
df['Average_Spending_Per_Purchase'] = pd.
    ↳to_numeric(df['Average_Spending_Per_Purchase'], errors='coerce')
df['Purchase_Frequency_Per_Month'] = pd.
    ↳to_numeric(df['Purchase_Frequency_Per_Month'], errors='coerce')
df['Brand_Affinity_Score'] = pd.to_numeric(df['Brand_Affinity_Score'],
    ↳errors='coerce')
df['Month'] = pd.to_numeric(df['Month'], errors='coerce')
```

```
df['Year'] = pd.to_numeric(df['Year'], errors='coerce')
```

```
[8]: # Select only numeric columns
numeric_columns = ['Age', 'Purchase_Amount', 'Average_Spending_Per_Purchase',
                   'Purchase_Frequency_Per_Month', 'Brand_Affinity_Score',
                   'Month', 'Year']

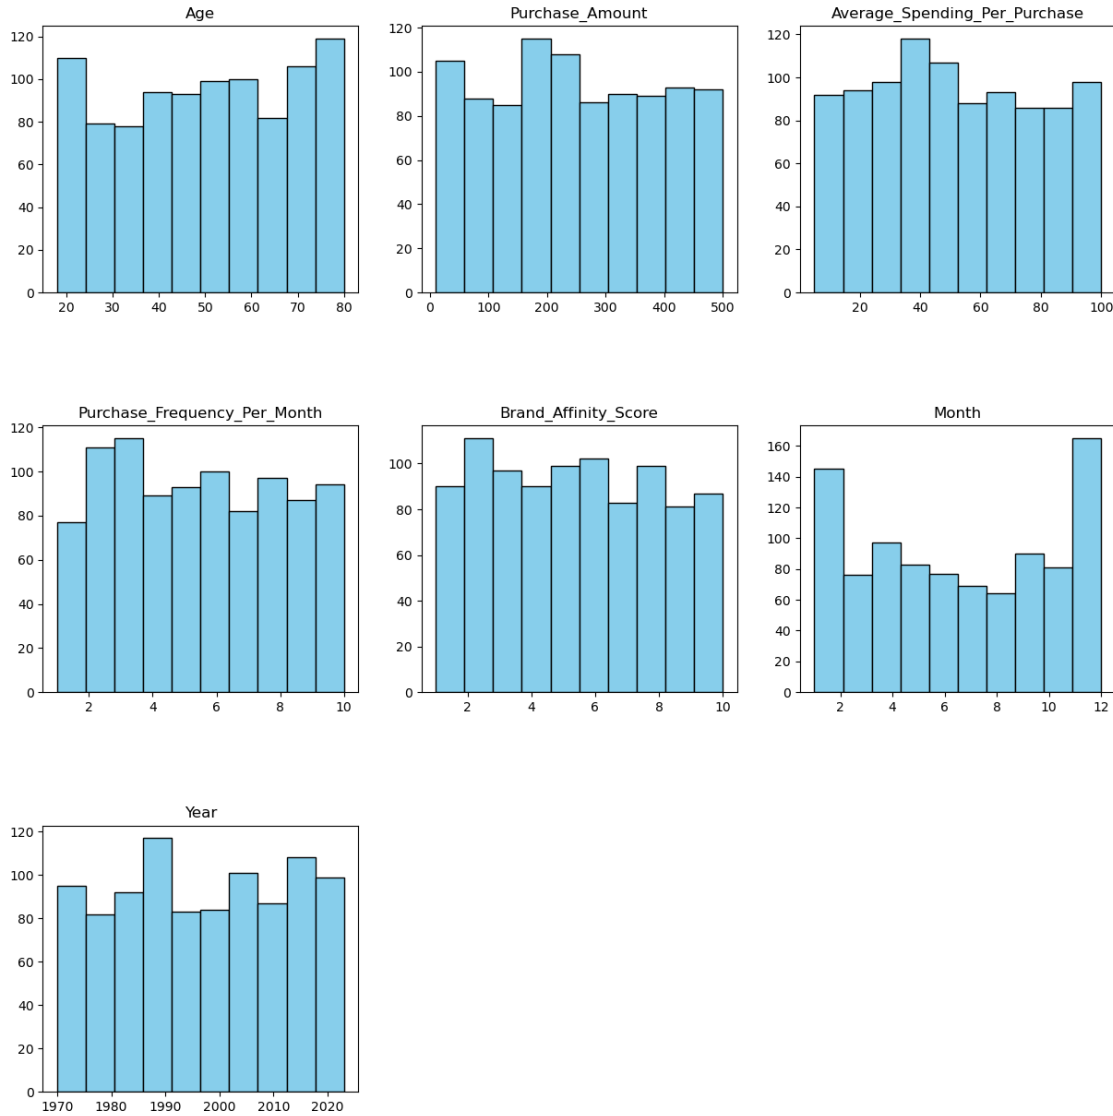
# Create subplots with 3 columns in each row
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(15, 15))
fig.subplots_adjust(hspace=0.5) # Adjust vertical spacing between subplots

# Flatten the 2D array of subplots to make it easier to iterate
axes = axes.flatten()

# Iterate through numeric columns and plot histograms
for i, column in enumerate(numeric_columns):
    df[column].plot(kind='hist', ax=axes[i], edgecolor='black', color='skyblue')
    axes[i].set_title(column)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('')

# Hide empty subplots
for j in range(len(numeric_columns), len(axes)):
    axes[j].axis('off')

plt.show()
```



```
[39]: # df.dropna(subset=['Customer_ID'], inplace=True)

df['Customer_ID'].fillna('Unknown', inplace=True)

df['Age'].fillna(df['Age'].median(), inplace=True)

df['Gender'].fillna(df['Gender'].mode()[0], inplace=True)

df['Income_Level'].fillna(df['Income_Level'].mode()[0], inplace=True)

df['Address'].fillna('Unknown', inplace=True)

df['Transaction_ID'].fillna('Unknown', inplace=True)
```

```

df['Purchase_Date'].fillna(method='ffill', inplace=True)

df['Product_ID'].fillna('Unknown', inplace=True)

df['Product_Category'].fillna(df['Product_Category'].mode()[0], inplace=True)

df['Brand'].fillna(df['Brand'].mode()[0], inplace=True)

df['Purchase_Amount'].fillna(df['Purchase_Amount'].median(), inplace=True)

df['Average_Spending_Per_Purchase'].fillna(df['Average_Spending_Per_Purchase'].
    ↪median(), inplace=True)

df['Purchase_Frequency_Per_Month'].fillna(df['Purchase_Frequency_Per_Month'].
    ↪median(), inplace=True)

df['Brand_Affinity_Score'].fillna(df['Brand_Affinity_Score'].median(),
    ↪inplace=True)

df['Product_Category_Preferences'].fillna(df['Product_Category_Preferences'].
    ↪mode()[0], inplace=True)

df['Month'].fillna(df['Month'].mode()[0], inplace=True)

df['Year'].fillna(df['Year'].mode()[0], inplace=True)

df['Season'].fillna(df['Season'].mode()[0], inplace=True)

df

```

```

[39]:
      Customer_ID  Age  Gender  Income_Level  \
0    b81ee6c9-2ae4-48a7-b283-220eaa244f43  40.0  Female      Medium
1                                Unknown  25.0    Male        High
2    fdf79bcd-5908-4c90-8501-570ffb5b7648  57.0   Other        Low
3    878dccba-893a-48f9-8d34-6ed394fa3c9c  38.0  Female      Medium
4    0af0bd81-73cc-494e-aa5e-75c6d0b6d743  68.0   Other      Medium
..          ...    ...    ...    ...
995          Unknown  70.0    Male      Medium
996  2116266d-8d1c-48cc-ac28-e4e675cb2a4d  78.0  Female        Low
997  562cee08-f909-4e1c-a811-5711f967bea5  63.0    Male      High
998  84da2eea-6e9e-46d4-8d94-1e9b0c377d78  43.0    Male      High
999  87629baf-a138-4374-be37-8bab776379b8  19.0   Other      High

      Address  \
0  43548 Murray Islands Suite 974\nAmyberg, CT 13457
1                                Unknown

```

2 79683 Kevin Hill Apt. 555\nJohnshire, AR 39961
3 02998 Hall Meadows Suite 809\nNorth Robertvill...
4 21411 Timothy Ford Apt. 320\nDavisborough, AR ...
..
995 566 Butler Turnpike\nPort Holly, OK 22329
996 45710 Wilson Circles Apt. 411\nWalterton, NC 8...
997 243 Emily Creek\nSouth Lindaport, CO 81594
998 1129 Kirby Ferry Suite 743\nBillyfurt, UT 41587
999 896 Troy Branch\nAmytown, NJ 62321

	Transaction_ID	Purchase_Date	\
0	c6a6c712-e36b-406a-bfde-f53bdcf4744f	2022-04-26	
1	0b587838-1e4f-4231-b488-42bcd47c052a	2021-08-10	
2	462925b1-a5bf-4996-bda2-59749de64eea	2021-12-09	
3	3cfafa02-6b34-4d77-9e05-d223dfab64e8	2022-12-03	
4	0d8dc27a-0c8f-4a82-b57e-8bf54cee9759	2020-06-08	
..	
995	776be313-5308-468e-a0ed-7409a4303364	2023-03-17	
996	51f771bf-2562-46c1-a25d-2f46f4bb1525	2023-08-30	
997	74eba598-ee91-4396-a137-6b869702ef29	2023-08-30	
998	4d2e213e-bcc0-4a8a-9501-6ca8361381c4	2021-05-13	
999	69afa592-2658-48ac-9b37-33a3a473d0be	2022-09-13	

	Product_ID	Product_Category	Brand	\
0	d2f767d6-b01a-41a2-87f7-ec1d1186f50e	Clothing	Brand_C	
1	79eadc55-2de1-41cf-b1b6-40118c0bf8ec	Books	Brand_A	
2	9ab75a68-4329-4bd9-a259-2233c0f34c93	Electronics	Brand_A	
3	d518569b-ff79-494b-b2b6-7e2af39db86a	Clothing	Brand_C	
4	b6deac9d-2b7e-4a51-8273-a6534910b3bc	Books	Brand_B	
..	
995	1802f115-80d8-48fd-ad97-94038fe31b82	Electronics	Brand_C	
996	546d8d8f-1498-4aa9-8123-29550d911a17	Books	Brand_B	
997	8b6ffec8-de54-445c-90d0-1399858b2e16	Electronics	Brand_C	
998	51ed2d86-c9ab-4922-a8ff-469acf6ac91e	Clothing	Brand_C	
999	91ba2109-15aa-40a0-aa9c-732a1e2e1e27	Clothing	Brand_B	

	Purchase_Amount	Average_Spending_Per_Purchase	\
0	193.0	59.0	
1	318.0	77.0	
2	197.0	100.0	
3	262.0	97.0	
4	429.0	85.0	
..	
995	180.0	92.0	
996	176.0	53.0	
997	212.0	99.0	
998	246.0	98.0	

999 208.0 12.0

	Purchase_Frequency_Per_Month	Brand_Affinity_Score \
0	2.0	2.0
1	2.0	1.0
2	9.0	1.0
3	3.0	4.0
4	7.0	2.0
..
995	2.0	5.0
996	3.0	3.0
997	2.0	9.0
998	8.0	7.0
999	10.0	1.0

	Product_Category_Preferences	Month	Year	Season
0	Low	1.0	2010.0	Winter
1	Low	8.0	1989.0	Fall
2	Low	4.0	1995.0	Winter
3	Low	9.0	2012.0	Fall
4	High	1.0	2010.0	Summer
..
995	Medium	5.0	1987.0	Fall
996	Medium	9.0	1977.0	Winter
997	Low	12.0	1995.0	Summer
998	Low	3.0	2000.0	Fall
999	Medium	12.0	1970.0	Summer

[1000 rows x 18 columns]

```
[40]: df.isna().sum()
```

```
[40]: Customer_ID      0
      Age              0
      Gender           0
      Income_Level     0
      Address          0
      Transaction_ID   0
      Purchase_Date    0
      Product_ID       0
      Product_Category 0
      Brand            0
      Purchase_Amount  0
      Average_Spending_Per_Purchase 0
      Purchase_Frequency_Per_Month  0
      Brand_Affinity_Score           0
      Product_Category_Preferences   0
```



```
Month          0
Year           0
Season         0
dtype: int64
```

```
[ ]:
```

```
[11]: # Select only numeric columns
numeric_columns = ['Age', 'Purchase_Amount', 'Average_Spending_Per_Purchase',
                   'Purchase_Frequency_Per_Month', 'Brand_Affinity_Score',
                   'Month', 'Year']

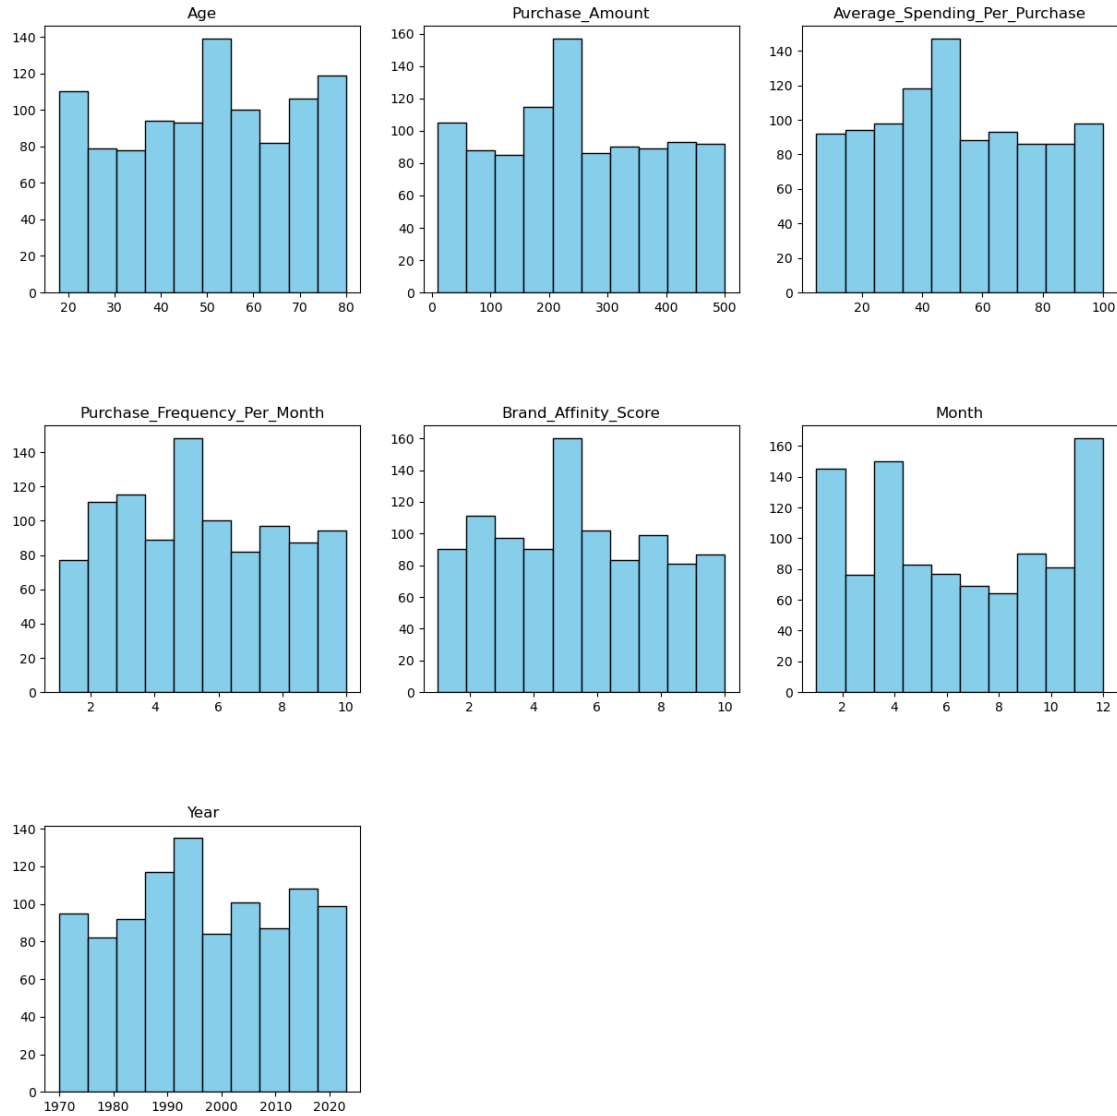
# Create subplots with 3 columns in each row
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(15, 15))
fig.subplots_adjust(hspace=0.5) # Adjust vertical spacing between subplots

# Flatten the 2D array of subplots to make it easier to iterate
axes = axes.flatten()

# Iterate through numeric columns and plot histograms
for i, column in enumerate(numeric_columns):
    df[column].plot(kind='hist', ax=axes[i], edgecolor='black', color='skyblue')
    axes[i].set_title(column)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('')

# Hide empty subplots
for j in range(len(numeric_columns), len(axes)):
    axes[j].axis('off')

plt.show()
```



```
[31]: import pandas as pd
from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Select relevant columns for clustering
columns_for_clustering = ['Purchase_Amount', 'Average_Spending_Per_Purchase']

# Convert selected columns to float, handling missing values
df[columns_for_clustering] = df[columns_for_clustering].apply(pd.to_numeric,
    errors='coerce')

# Drop rows with missing values
```

```

df.dropna(subset=columns_for_clustering, inplace=True)

# Standardize the data
scaler = StandardScaler()
data_for_clustering = scaler.fit_transform(df[columns_for_clustering])

# Apply DBSCAN
dbscan = DBSCAN(eps=0.5, min_samples=5)
df['Cluster'] = dbscan.fit_predict(data_for_clustering)

# Get cluster labels
labels = df['Cluster'].unique()

# Plotting
plt.figure(figsize=(10, 6))


# Plot each cluster separately
for label in labels:
    cluster_points = df[df['Cluster'] == label]

    if label == -1:
        # Border points (outliers)
        plt.scatter(cluster_points['Purchase_Amount'],
            ↪cluster_points['Average_Spending_Per_Purchase'],
                    c='black', label='Border Points', marker='x', s=100)
    else:
        # Core points
        plt.scatter(cluster_points['Purchase_Amount'],
            ↪cluster_points['Average_Spending_Per_Purchase'],
                    label=f'Cluster {label}', s=50, alpha=0.7)

# Centroids of each cluster
centroids = df.groupby('Cluster')[columns_for_clustering].mean()
plt.scatter(centroids['Purchase_Amount'],
    ↪centroids['Average_Spending_Per_Purchase'],
            marker='o', color='red', s=200, edgecolors='black', linewidths=2,
    ↪label='Centroids')

plt.title('DBSCAN Clustering')
plt.xlabel('Purchase Amount')
plt.ylabel('Average Spending Per Purchase')
plt.legend()
plt.show()

```



```
[44]: import pandas as pd
from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import numpy as np

# Extract the numerical columns for clustering
numerical_columns = ['Average_Spending_Per_Purchase', 'Purchase_Frequency_Per_Month']
X = df[numerical_columns]

# Standardize the numerical data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply DBSCAN
dbscan = DBSCAN(eps=0.5, min_samples=5)
labels = dbscan.fit_predict(X_scaled)

# Add the cluster labels to the original DataFrame
df['Cluster_Labels'] = labels

# Visualize the results (scatter plot for two numerical columns)
plt.figure(figsize=(10, 6))
```

```

# Define colors for each cluster (including noise points)
unique_labels = np.unique(labels)
colors = plt.cm.rainbow(np.linspace(0, 1, len(unique_labels)))

# Plot each cluster with different markers
for cluster_label, color in zip(unique_labels, colors):
    cluster_mask = (labels == cluster_label)

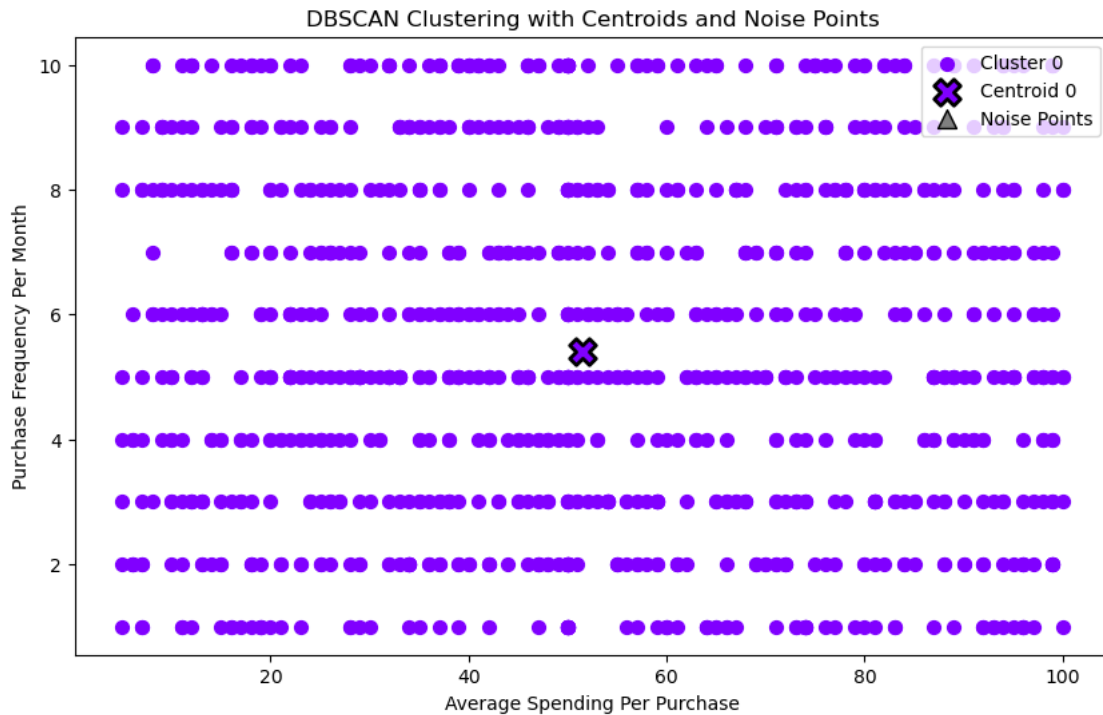
    # Plot cluster points
    plt.scatter(X[cluster_mask]['Average_Spending_Per_Purchase'],
                X[cluster_mask]['Purchase_Frequency_Per_Month'],
                c=[color], label=f'Cluster {cluster_label}', marker='o', s=50)

    # Highlight cluster centroid (mean of cluster points)
    if cluster_label != -1: # Avoid calculating centroid for noise points
        centroid = X[cluster_mask].mean()
        plt.scatter(centroid['Average_Spending_Per_Purchase'],
                    centroid['Purchase_Frequency_Per_Month'],
                    c=[color], marker='X', s=200, edgecolors='k', linewidth=2,
                    label=f'Centroid {cluster_label}')

# Identify noise points (points not belonging to any cluster)
noise_points_mask = (labels == -1)
plt.scatter(X[noise_points_mask]['Average_Spending_Per_Purchase'],
            X[noise_points_mask]['Purchase_Frequency_Per_Month'],
            c='gray', marker='^', s=100, edgecolors='k', linewidth=1,
            label='Noise Points')

plt.title('DBSCAN Clustering with Centroids and Noise Points')
plt.xlabel('Average Spending Per Purchase')
plt.ylabel('Purchase Frequency Per Month')
plt.legend()
plt.show()

```



```
[58]: import pandas as pd
import numpy as np
from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler, LabelEncoder
import matplotlib.pyplot as plt

# Load the data

# Select numerical columns for clustering
numerical_cols = ['Average_Spending_Per_Purchase', 'Age']
X = df[numerical_cols]

# Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply DBSCAN
dbscan = DBSCAN(eps=0.5, min_samples=5)
labels = dbscan.fit_predict(X_scaled)

# Visualize the results
plt.figure(figsize=(10, 6))

# Scatter plot for clusters
```

```

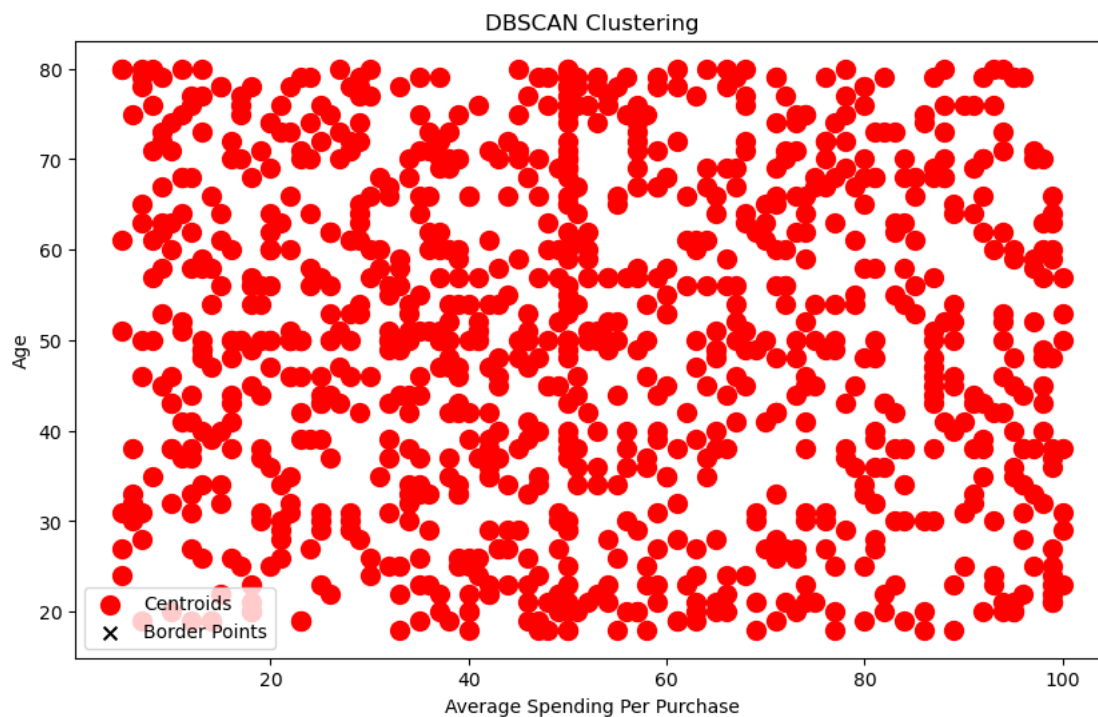
scatter = plt.scatter(X['Average_Spending_Per_Purchase'], X['Age'], c=labels,
    cmap='viridis', s=50, alpha=0.8)

# Plot centroids (core samples)
core_samples_mask = np.zeros_like(labels, dtype=bool)
core_samples_mask[dbscan.core_sample_indices_] = True
core_samples = X[core_samples_mask]
plt.scatter(core_samples['Average_Spending_Per_Purchase'], core_samples['Age'],
    c='red', marker='o', s=100, label='Centroids')

# Plot border points
border_points = X[labels == -1]
plt.scatter(border_points['Average_Spending_Per_Purchase'],
    border_points['Age'], c='black', marker='x', s=50, label='Border Points')

plt.title('DBSCAN Clustering')
plt.xlabel('Average Spending Per Purchase')
plt.ylabel('Age')
plt.legend()
plt.show()

```



```

[61]: import pandas as pd
import numpy as np

```

```

from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Extract the relevant numerical columns
numerical_columns = ['Age', 'Purchase_Frequency_Per_Month']
X = df[numerical_columns]

# Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply DBSCAN
dbscan = DBSCAN(eps=0.5, min_samples=5)
labels = dbscan.fit_predict(X_scaled)

# Add cluster labels to the DataFrame
df['Cluster'] = labels

# Plot the data points, centroids, and border points
core_samples_mask = np.zeros_like(dbscan.labels_, dtype=bool)
core_samples_mask[dbscan.core_sample_indices_] = True

unique_labels = set(labels)
colors = [plt.cm.Spectral(each) for each in np.linspace(0, 1, len(unique_labels))]

plt.figure(figsize=(10, 6))

# Plot core samples in black
plt.scatter(X.iloc[core_samples_mask, 0], X.iloc[core_samples_mask, 1], c='k',
            marker='o', s=100, label='Core Samples')

for k, col in zip(unique_labels, colors):
    if k == -1:
        col = [0, 0, 0, 1]

    class_member_mask = (labels == k)


    xy = X[class_member_mask & core_samples_mask]
    plt.scatter(xy.iloc[:, 0], xy.iloc[:, 1], c=[col], marker='o', s=50,
                label=f'Cluster {k} (Centroid)')

    xy = X[class_member_mask & ~core_samples_mask]
    plt.scatter(xy.iloc[:, 0], xy.iloc[:, 1], c=[col], marker='o', s=30,
                label=f'Cluster {k} (Border Point)')

```



```
plt.title('DBSCAN Clustering with Core Samples, Centroids, and Border Points')
plt.xlabel('Age')
plt.ylabel('Purchase Frequency Per Month')
plt.legend()
plt.show()
```



```
[66]: import pandas as pd
import numpy as np
from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Open the data file

# Extract relevant columns
data = df[['Purchase_Amount', 'Month']]

# Standardize the data
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data)

# Apply DBSCAN
dbscan = DBSCAN(eps=0.5, min_samples=5)
labels = dbscan.fit_predict(data_scaled)
```

```


# Visualize the results
plt.figure(figsize=(10, 6))

# Core points
core_samples_mask = np.zeros_like(labels, dtype=bool)
core_samples_mask[dbscan.core_sample_indices_] = True
plt.scatter(data['Purchase_Amount'][core_samples_mask],
            data['Month'][core_samples_mask],
            c=labels[core_samples_mask], cmap='viridis', marker='o', s=50,
            label='Core Points')

# Border points
border_points_mask = (labels == -1) & ~core_samples_mask
plt.scatter(data['Purchase_Amount'][border_points_mask],
            data['Month'][border_points_mask],
            c='red', marker='x', s=50, label='Border Points')

plt.title('DBSCAN Clustering with Border Points')
plt.xlabel('Purchase Amount')
plt.ylabel('Month')
plt.legend()
plt.show()

```



[]: